

KUBERNETES & TERRAFORM ARCHITECTURE

Complete Infrastructure Guide for DataMigrate AI

Author: Alexander Garcia Angus
Property of: OKO Investments

Generated: December 02, 2025

1. COMPLETE TECHNOLOGY STACK

Infrastructure Layer (Terraform):

- * Amazon EKS - Kubernetes orchestration
- * Amazon VPC - Network isolation
- * Amazon RDS PostgreSQL - Relational database
- * Amazon ElastiCache Redis - Caching & message broker
- * Amazon ECR - Docker image registry
- * AWS Secrets Manager - Secure credential storage
- * Amazon S3 - Object storage & frontend hosting
- * Amazon CloudFront - CDN for frontend
- * AWS ALB - Application Load Balancer
- * AWS CloudWatch - Monitoring & logging

Application Layer (Kubernetes):

- * Go Backend - RESTful API (Gin framework)
- * Vue.js 3 Frontend - Modern SPA (TypeScript)
- * Python AI Service - LangGraph Agents
- * Celery Workers - Background task processing
- * Nginx Ingress - Kubernetes ingress controller

2. ARCHITECTURE OVERVIEW

The architecture follows a modern cloud-native design with the following flow:

| Layer | Components | Purpose |
|----------------|-----------------------------|--|
| DNS/CDN | Route53, CloudFront, WAF | DNS routing, edge caching, security |
| Load Balancing | AWS ALB | HTTPS termination, traffic distribution |
| Orchestration | Amazon EKS | Container orchestration (K8s 1.28) |
| Application | Go API, Python AI, Celery | Business logic, AI agents, background jobs |
| Data | RDS PostgreSQL, ElastiCache | Persistence, caching, message broker |
| Storage | S3, ECR | Artifacts, Docker images |
| Security | Secrets Manager, IAM IRSA | Credentials, service accounts |

3. REQUEST FLOWS

Frontend Request (Vue.js):

User Browser -> CloudFront CDN (edge cache) -> S3 Static Website -> Downloads index.html, app.js, app.css
-> User sees Vue.js app

API Request (Go Backend):

Vue.js App -> API call (fetch/axios) -> ALB -> Nginx Ingress -> Go API Pod (1 of 3 replicas) -> Redis cache check -> PostgreSQL if miss -> Response

Migration Request (AI Agents):

User initiates migration -> Go API receives request -> Creates Celery task -> Redis queue -> Celery Worker -> Python AI Service -> LangGraph multi-agent workflow:

1. Metadata Extraction Agent
2. Schema Analysis Agent
3. dbt Model Generator Agent
4. Validator Agent
5. Orchestrator Agent

-> Saves results to PostgreSQL -> Updates status -> Frontend polls for updates

4. KUBERNETES DEPLOYMENTS

| Service | Replicas | Resources | Auto-scaling |
|-------------------|----------|---------------------|---------------------|
| Go API | 3 | 512Mi RAM, 500m CPU | 3-20 pods (70% CPU) |
| Python AI Service | 2 | 1Gi RAM, 1000m CPU | 2-10 pods |
| Celery Workers | 5 | 1Gi RAM, 1000m CPU | 2-20 pods (75% CPU) |
| Nginx Ingress | 2 | 256Mi RAM, 250m CPU | Fixed |

Key Kubernetes Features:

- * Health checks (liveness + readiness probes)
- * Horizontal Pod Autoscaler (HPA) based on CPU/memory
- * External Secrets Operator for AWS Secrets Manager integration
- * TLS termination with cert-manager (Let's Encrypt)
- * Resource limits and requests for QoS

5. COST BREAKDOWN

Development Environment (~\$250/month):

| Service | Configuration | Monthly Cost |
|---------------------|----------------|--------------------|
| EKS Control Plane | 1 cluster | \$73.00 |
| EKS Worker Nodes | 2x t3.medium | \$60.00 |
| RDS PostgreSQL | db.t3.micro | \$14.00 |
| ElastiCache Redis | cache.t3.micro | \$12.00 |
| NAT Gateways | 3x NAT | \$32.40 |
| ALB | Application LB | \$16.00 |
| S3 + CloudFront | Frontend | \$10.00 |
| Other (ECR, CW, SM) | Various | \$40.00 |
| TOTAL | | \$258/month |

Production Environment (~\$1,200-2,000/month):

| Service | Configuration | Monthly Cost |
|-------------------|------------------------|--------------|
| EKS Control Plane | 1 cluster | \$73.00 |
| EKS Worker Nodes | 4-10x t3.large (avg 6) | \$375.00 |

| | | |
|-------------------|--------------------------|----------------------|
| RDS PostgreSQL | db.t3.large Multi-AZ | \$280.00 |
| ElastiCache Redis | cache.m5.large (3 nodes) | \$260.00 |
| NAT + ALB + S3/CF | Production config | \$150.00 |
| Monitoring + WAF | CloudWatch + WAF | \$130.00 |
| TOTAL | | \$1,331/month |

Cost Optimization Tips:

- * Use Reserved Instances for RDS (40% savings)
- * Use Spot Instances via Karpenter (70% savings)
- * Implement auto-scaling (scale down during off-hours)
- * Use VPC Endpoints to eliminate NAT costs

6. DEPLOYMENT WORKFLOW

Step 1: Build Docker Images

Build Go API, Python AI Service, and Celery Worker images. Tag and push to ECR.

Step 2: Deploy Infrastructure (Terraform)

```
cd terraform/environments/dev && terraform init && terraform plan && terraform apply
```

Step 3: Configure kubectl

```
aws eks update-kubeconfig --region us-east-1 --name datamigrate-ai-dev-eks
```

Step 4: Deploy Kubernetes Resources

```
kubectl create namespace datamigrate-ai && kubectl apply -f k8s/external-secrets.yaml && kubectl apply -f k8s/deployments/ && kubectl apply -f k8s/ingress.yaml
```

Step 5: Deploy Frontend

```
cd frontend && npm run build && aws s3 sync dist/ s3://datamigrate-ai-frontend --delete && aws cloudfront create-invalidation --distribution-id --paths '/'
```

7. TECHNOLOGY STACK SUMMARY

| Category | Technology | Purpose |
|---------------|-----------------------|-------------------------|
| IaC | Terraform | Infrastructure as Code |
| Orchestration | Kubernetes (EKS) | Container orchestration |
| Backend | Go (Gin) + Python | REST API + AI agents |
| Frontend | Vue.js 3 + TypeScript | SPA framework |
| AI Framework | LangGraph + LangChain | Multi-agent system |
| Task Queue | Celery + Redis | Background jobs |
| Database | PostgreSQL 15 (RDS) | Relational data |
| Cache | Redis (ElastiCache) | Caching + broker |
| CDN | CloudFront | Frontend delivery |
| CI/CD | GitHub Actions | Automated deployment |

This is an enterprise-grade, production-ready architecture designed for OKO Investments!

Copyright 2025 OKO Investments. All rights reserved.