

# DataMigrate AI

## Interview Preparation Guide

Alexander Garcia Angus

OKO Investments

**Project:** MSSQL to dbt Migration SaaS Platform powered by AI agents

## 1. PROJECT OVERVIEW

DataMigrate AI is an AI-powered SaaS platform that automates the migration of legacy MSSQL databases to modern dbt (data build tool) projects. The platform uses a multi-agent system built with LangGraph and Claude API to analyze database schemas, generate dbt models, and validate the transformations.

## Technology Stack

Component	Technology	Purpose
Frontend	Vue.js 3 + TypeScript	Modern reactive UI with type safety
Backend API	FastAPI + Python 3.12	Async REST API with auto-docs
AI Agents	LangGraph + Claude API	Multi-agent workflow orchestration
Database	PostgreSQL + Redis	Persistent storage + caching
Infrastructure	Kubernetes (EKS)	Container orchestration
IaC	Terraform	Infrastructure as Code
Autoscaling	Karpenter	40-60% cost savings

## 2. ARCHITECTURE INTERVIEW QUESTIONS

### Q1: Explain the overall architecture of DataMigrate AI

**Answer:** DataMigrate AI uses a 3-tier architecture:

- 1. Frontend (Vue.js 3 + TypeScript):** Single-page application with Pinia state management. Users authenticate, create migrations, and monitor progress in real-time with auto-refresh every 10-30 seconds.
- 2. Backend (FastAPI):** Async REST API that handles authentication, migration CRUD operations, and orchestrates LangGraph agents. Uses PostgreSQL for persistent state and Redis for caching.
- 3. AI Layer (LangGraph + Claude):** Multi-agent system with 6 specialized agents: Assessment, Planner, Executor, Tester, Rebuilder, and Evaluator. Each agent has a specific role in the migration workflow.

The frontend communicates with FastAPI via REST, FastAPI orchestrates LangGraph agents, and agents persist their state to PostgreSQL for checkpoint recovery.

### Q2: Why did you choose LangGraph over LangChain?

**Answer:** I chose LangGraph because migrations can take 30+ minutes and require stateful execution:

- 1. State Persistence:** LangGraph provides built-in checkpointing to PostgreSQL. If a migration fails, it can resume from the last checkpoint instead of starting over.
- 2. Agent Communication:** LangGraph has native support for inter-agent message passing, which LangChain lacks.
- 3. Complex Workflows:** LangGraph's graph-based execution allows conditional routing (e.g., if validation fails, route to rebuilder agent).

**Trade-off:** LangGraph has a steeper learning curve than LangChain, but the benefits for long-running workflows justify it. For simple Q&A; agents, LangChain would be sufficient.

### Q3: Why Kubernetes (EKS) instead of ECS Fargate?

**Answer:** I chose Kubernetes for three strategic reasons:

- 1. Cost Optimization with Karpenter:** Karpenter provides intelligent node autoscaling that saves 40-60% on compute costs compared to standard Cluster Autoscaler. Analysis showed \$960-4,200/year savings.
- 2. Skill Transferability:** Kubernetes skills are cloud-agnostic. If OKO Investments expands to Azure or GCP, the same K8s knowledge applies. ECS is AWS-only.
- 3. Mature Ecosystem:** Kubernetes has better horizontal pod autoscaling (HPA), more third-party tools, and stronger community support.

**Trade-off:** EKS control plane costs \$73/month vs ECS Fargate's \$0 control plane. But Karpenter savings offset this within 1 month.

### 3. TECHNICAL IMPLEMENTATION QUESTIONS

#### Q4: How do you handle failures in long-running migrations?

**Answer:** I implemented a checkpoint system with three layers of resilience:

**1. Agent-Level Checkpoints:** Each LangGraph agent saves its state after completing a step. State is persisted to PostgreSQL with the migration\_id as the key.

**2. Table-Level Progress:** The migration tracks which tables have been processed. If interrupted, it resumes from the next unprocessed table, not from the beginning.

**3. Retry Logic:** Failed tables are retried up to 3 times with exponential backoff (1s, 2s, 4s). If all retries fail, the table is marked as 'failed' but the migration continues with other tables.

**Testing:** I validated this with chaos engineering - manually killing FastAPI workers and corrupting checkpoint data to ensure graceful degradation.

#### Q5: Explain your authentication strategy

**Answer:** DataMigrate AI uses two authentication mechanisms:

**1. Session-Based (Frontend):** Users log in via Vue.js frontend. FastAPI returns a JWT token stored in httpOnly cookies (prevents XSS). The token expires after 24 hours.

**2. API Key-Based (External API):** External services use API keys (format: mk\_xxxxx). Keys are stored hashed in PostgreSQL using bcrypt. Each key has a rate limit (default: 100 req/hour).

##### Security Considerations:

- Passwords hashed with bcrypt (cost factor 12)
- JWT signed with HS256 + secret key from environment variables
- API keys never logged (redacted in logs)
- Rate limiting prevents abuse

#### Q6: How would you optimize a slow SQL query?

**Answer:** My systematic approach:

**1. Identify the Problem:** Use EXPLAIN ANALYZE to see the query plan. Look for sequential scans, high cost estimates, or slow actual times.

##### 2. Common Fixes:

- Add indexes on WHERE/JOIN columns
- Fix N+1 queries with eager loading (SQLAlchemy joinedload())
- Avoid SELECT \* (specify only needed columns)
- Use query result caching for repeated reads

**Real Example from DataMigrate AI:** Migration queries were slow because I was loading related model\_files in a loop (N+1 problem). I switched to:

```
query = select(Migration).options(joinedload(Migration.model_files))
```

This reduced 100 queries to 1, cutting load time from 5s to 200ms.

## 4. FRONTEND & USER EXPERIENCE QUESTIONS

### Q7: Explain your Vue.js state management strategy

**Answer:** I use Pinia for centralized state management with two main stores:

#### 1. Auth Store (auth.ts):

- Manages user authentication state (user, token, isAuthenticated)
- Provides login(), logout(), fetchCurrentUser() actions
- Persists token to localStorage for page refresh

#### 2. Migrations Store (migrations.ts):

- Manages migration list and current migration state
- Provides fetchMigrations(), createMigration(), deleteMigration() actions
- Implements polling for real-time updates

**Why Pinia over Vuex?** Pinia has better TypeScript support, simpler API (no mutations), and is the official recommendation for Vue 3.

### Q8: How do you handle real-time updates in the UI?

**Answer:** I implemented polling with exponential backoff:

**Dashboard:** Polls every 30 seconds for migration statistics. Longer interval because stats change slowly.

**Migrations List:** Polls every 10 seconds to show updated status and progress.

**Why Polling vs WebSockets?** For MVP, polling is simpler and sufficient. WebSockets would add complexity (connection management, reconnection logic) for minimal benefit at current scale. I'd switch to WebSockets at 1000+ concurrent users for lower server load.

**Optimization:** Polling stops when migration completes (status === 'completed'). This prevents unnecessary API calls.

## 5. INFRASTRUCTURE & DEVOPS QUESTIONS

### Q9: Explain your Terraform infrastructure setup

**Answer:** I organized Terraform into reusable modules:

#### **Module Structure:**

- vpc/ - 3 AZ VPC with public, private, and database subnets
- eks/ - Kubernetes cluster with Karpenter autoscaler
- rds/ - PostgreSQL Multi-AZ with automated backups
- redis/ - ElastiCache for caching and Celery broker
- s3/ - Storage for dbt models and logs

#### **Key Decisions:**

1. Multi-AZ deployment for 99.9% uptime SLA
2. Private subnets for EKS (no direct internet access)
3. NAT Gateway for outbound traffic (ECR image pulls)
4. RDS encryption at rest with KMS

**Cost Optimization:** Dev environment uses t3.small (1 node), Production uses m5.large (3+ nodes with autoscaling).

### Q10: What monitoring and observability have you implemented?

**Answer:** Three-layer observability strategy:

#### **1. Application Metrics (CloudWatch):**

- API latency (p50, p95, p99)
- Migration success/failure rates
- Database connection pool usage
- Custom metric: Time per migration phase

#### **2. Infrastructure Metrics (Kubernetes):**

- Pod CPU/memory usage
- Node autoscaling events (Karpenter)
- PVC (storage) usage

#### **3. Logging (CloudWatch Logs):**

- Structured JSON logs with correlation IDs
- Log levels: ERROR for failures, INFO for events
- Logs retained for 30 days (compliance)

#### **Alerting:** SNS alerts for:

- API error rate > 5%
- Database CPU > 80%
- Migration failures

## 6. PERFORMANCE & SCALABILITY QUESTIONS

### Q11: How does your system scale to handle increased load?

**Answer:** Multi-layer scaling strategy:

#### 1. Application Layer (Kubernetes HPA):

- FastAPI pods scale 2-20 based on CPU > 70%
- Celery worker pods scale 1-10 based on queue depth
- Average scale-up time: 30 seconds

#### 2. Infrastructure Layer (Karpenter):

- Adds EC2 nodes when pod scheduling fails
- Consolidates nodes when underutilized (cost savings)
- Uses spot instances for 70% cost reduction

#### 3. Database Layer:

- RDS read replicas for read-heavy queries
- Connection pooling (max 100 connections)
- Redis caching for frequently accessed data (TTL: 5 minutes)

**Bottleneck Monitoring:** CloudWatch alarms trigger when scaling is needed. Manual intervention for database vertical scaling (larger instance).

### Q12: When would you add Rust microservices?

**Answer:** I'd add Rust when specific bottlenecks emerge:

#### Trigger Conditions:

1. API latency p95 > 1 second consistently
2. CPU usage > 80% with optimized Python code
3. Processing > 10,000 migrations/month

#### Candidate Services for Rust:

1. SQL Parser (5s → 500ms, 10x faster)
2. dbt Compiler (10s → 1s, 10x faster)
3. Schema Validator (2s → 200ms, 10x faster)

**Hybrid Strategy:** Keep FastAPI for 80% (CRUD, auth, orchestration), add Rust for 20% bottlenecks. FastAPI calls Rust microservices via HTTP.

**Expected Benefits:** 50-70% cost reduction, 10x performance improvement, better user experience (faster migrations).

## 7. TESTING & QUALITY ASSURANCE QUESTIONS

### Q13: What testing strategy did you implement?

**Answer:** Multi-layer testing pyramid:

#### 1. Unit Tests (pytest):

- Test individual functions (SQL parsing, dbt generation)
- Mock external dependencies (Claude API, database)
- Coverage target: 80%

#### 2. Integration Tests:

- Test LangGraph agent workflow end-to-end
- Use test database (SQLite in CI, PostgreSQL locally)
- Verify checkpoint recovery

#### 3. E2E Tests (Playwright):

- Test critical user flows (login, create migration, view results)
- Run against staging environment before production deploy

**Chaos Testing:** Manually kill workers, corrupt checkpoints, simulate network failures to validate resilience.

### Q14: How do you ensure code quality?

**Answer:** Automated quality gates in CI/CD:

#### 1. Linting (Pre-commit hooks):

- Python: black (formatting), flake8 (linting), mypy (type checking)
- TypeScript: ESLint + Prettier
- Pre-commit hooks prevent committing bad code

#### 2. Security Scanning:

- Dependabot for dependency vulnerabilities
- Bandit for Python security issues
- Trivy for Docker image scanning

#### 3. Code Reviews:

- All changes require review before merge
- Checklist: Tests pass, documentation updated, no secrets committed

#### 4. CI Pipeline:

- Lint → Test → Build → Deploy
- Blocks merge if any step fails

## 8. BEHAVIORAL & PROJECT MANAGEMENT QUESTIONS

### Q15: Describe a technical challenge you faced and how you solved it

**Challenge:** LangGraph agents were failing randomly with 'state not found' errors.

**Investigation:**

1. Added detailed logging to trace execution flow
2. Discovered race condition: Multiple agents writing to same state key
3. PostgreSQL row locking wasn't sufficient due to async execution

**Solution:**

1. Implemented Redis distributed lock (SETNX command)
2. Each agent acquires lock before updating state
3. Lock expires after 30 seconds (prevents deadlock)

**Result:** Zero 'state not found' errors after fix. Migration success rate improved from 85% to 100%.

**Learning:** Async systems need explicit coordination. Row-level locks aren't enough for distributed systems.

### Q16: How do you prioritize features when building a product?

**Answer:** I use the RICE framework (Reach, Impact, Confidence, Effort):

**Example from DataMigrate AI:**

**High Priority (Built for MVP):**

- User authentication (High reach, high impact, high confidence, low effort)
- Basic migration workflow (High reach, critical impact)
- Real-time progress updates (High impact on UX)

**Medium Priority (Post-MVP):**

- Email notifications (Medium reach, medium impact)
- Advanced filtering (Medium impact, low effort)

**Low Priority (Future):**

- Rust microservices (Low reach initially, high effort)
- Multi-region deployment (Low reach for beta)

**Principle:** Deliver core value first, then optimize.

## 9. COST OPTIMIZATION & BUSINESS QUESTIONS

**Q17: How much does it cost to run DataMigrate AI monthly?**

**Answer:** Cost breakdown by environment:

Service	Dev (Monthly)	Production (Monthly)
EKS Control Plane	\$73	\$73
EC2 Nodes (Karpenter)	\$45 (t3.small)	\$180-360 (m5.large spot)
RDS PostgreSQL	\$25 (t4g.small)	\$120 (r6g.large Multi-AZ)
ElastiCache Redis	\$15 (t3.micro)	\$60 (r6g.large)
S3 Storage	\$5	\$20
CloudWatch/Logs	\$10	\$40
Data Transfer	\$5	\$30
<b>TOTAL</b>	<b>\$178/month</b>	<b>\$523-703/month</b>

**Cost Optimizations Implemented:**

1. Karpenter autoscaling saves 40-60% vs standard EC2
2. Spot instances for non-critical workloads (70% savings)
3. Redis caching reduces database queries by 60%
4. S3 Intelligent-Tiering for automatic archival

**Revenue Model:** \$49/month per user (breakeven at 15 users in production).

## 10. YOUR 60-SECOND ELEVATOR PITCH

**Practice this opening statement for interviews:**

"I'm Alexander Garcia Angus, and I built DataMigrate AI - an AI-powered SaaS platform that automates MSSQL to dbt migrations. The platform uses a multi-agent system with LangGraph and Claude API, where 6 specialized agents coordinate to analyze schemas, generate dbt models, and validate transformations.

The frontend is built with Vue.js 3 and TypeScript, providing real-time progress updates. The backend uses FastAPI with PostgreSQL for state management. Infrastructure runs on Kubernetes (EKS) with Karpenter autoscaling, which saves 40-60% on compute costs compared to standard autoscaling.

I implemented a checkpoint system that persists agent state every 30 seconds, allowing migrations to resume after interruptions - critical for using spot instances. The system has achieved 100% success rate on test migrations and is ready for production deployment.

I can walk you through the architecture, explain any technical decisions, or demonstrate the system live. What aspect would you like to dive deeper into?"

## 11. QUESTIONS TO ASK YOUR INTERVIEWER

Always prepare thoughtful questions. Here are strong technical questions:

- What's your current deployment strategy? (Shows interest in their infrastructure)
- How do you handle database migrations in production? (Technical depth)
- What observability tools do you use for distributed systems? (Shows ops awareness)
- What's your approach to incident response and on-call? (Work-life balance concern)
- How do you balance technical debt with new features? (Product thinking)
- What's your code review process? (Team collaboration)
- How do you make build vs buy decisions? (Business acumen)
- What's the most interesting technical challenge your team is facing? (Shows engagement)

## 12. QUICK REFERENCE CARD

Memorize these key facts:

Metric	Value
Migration Success Rate	100% (7/7 test models)
Agent Count	6 specialized agents
Average Migration Time	5-30 minutes (depends on size)
Checkpoint Frequency	Every 30 seconds
Cost Savings (Karpenter)	40-60% vs standard autoscaling
Database	PostgreSQL (Multi-AZ)
Cache	Redis (ElastiCache)
Frontend	Vue.js 3 + TypeScript
Backend	FastAPI + Python 3.12
Infrastructure	Kubernetes (EKS) + Terraform
Autoscaling	Karpenter (2.5min scale-up)
API Latency Target	p95 < 500ms
Uptime SLA	99.9% (Multi-AZ)

**Remember:** Confidence comes from genuine understanding. You built this system, you understand the trade-offs, and you can explain it clearly. Good luck!